Athlete Load and Wellness Report: Extensive Report

Perry Battles

May 2025

1 Key Insights and Action Items

1.1 Key Insights and Action Items

- EDA demonstrated that average session heart rate and practice load likely express much of the same information. If collecting both is logistically intractable or difficult, not much is lost by sacrificing one. Because of the familiarity of heart rate data and zones, it may be more useful to present these values to coaches and athletes.

- Sleep has by far the most outsized effect in promoting athlete wellness by reducing fatigue and stress and increasing mood and motivation (see figures 3 through 18).

- Workload does not appear to affect sleep quality within the ranges in which it has been administered (figure 19), with the exception of a potential slight effect for practice load.

- Soreness was negatively associated with sleep to a much greater degree than with any other variable (see 21). Competition load may positively influence soreness.

- There was an association between greater strength and conditioning load and lower levels of fatigue and stress (figures 3, **??**). However, this effect is likely very slight and possibly confounded by position-based factors; for example, it may be the case that positions that engage in more strength and conditioning perform less other training.

- To promote athlete well-being, facilitating proper restful sleep should be considered one of the highest priorities, especially during circumstances of travel and international competition. Very secondarily, the inclusion of resistance training may help to reduce levels of fatigue and stress. Finally, since soreness is associated with high competition load, techniques to mitigate this trend may be beneficial for performance. For example, during periods of dense competition frequency when acute performance is prioritized over chronic adaptation, utilizing recovery modalities that reduce soreness may enhance performance.

2 Introduction

The provided data outline day-by-day workload and wellness variables for each of 25 unique athletes on 334 different days in 2022. This report is aimed at the sports scientist, and discusses insights gained both from the exploration of this dataset and from subsequent modeling. Additionally, it outlines several action items derived from these insights designed to improve athlete performance and wellbeing by capitalizing upon the relationships discovered in the data.

The Methods and Results sections below are divided more or less into two portions: the first of each addresses exploratory data analysis, and the second touches on generalized linear modeling. The Methods section also includes a section on data preprocessing, which describes exactly how the data were aggregated for the purposes of further analysis. The Results section enumerates the patterns discovered in the data and the outcomes produced by fitting gamma regressions, and is followed by a summary of the findings, given in the Conclusion section, which restates the insights and action items provided above and outlines directions for future work.

3   Methods

3.1   Exploratory Data Analysis

Data were assessed for sparsity (e.g., number of missing values per feature), and distribution for the purposes of discovering patterns and guiding future modeling techniques. Additionally, multicollinearity was assessed to ascertain which variables could be included in the same set of predictors during modeling without undermining our ability to draw conclusions concerning the effect of one variable on the outcome of interest compared to another.

3.2   Data Preprocessing

Prior experience with workload quantification indicated that trends often emerge at the weekly and monthly levels which are absent when data are binned by individual calendar day. Hence, athlete data was grouped into weekly increments, with some variables (e.g., sleep) being summed for each day in the week and others (e.g., motivation) being averaged. Once these values had been collected, change in each value from the prior time window was also calculated so that subsequent modeling could determine if certain variables might affect the absolute value of one outcome but not the change therein or vice versa.

Prior to model fitting, the data were split into training and testing sets using a 70/30 train/test split. The data were standardized using a robust scaler, as this does not rely upon the same statistical assumptions as a standard scaler. This is valuable given that not all of the variables in the dataset are normally distributed (see figure 1). Robust scaling frees the analysis from this assumption.

3.3   Model Fitting

Since the outcome variables in question appeared to be roughly gamma-distributed, we selected a gamma generalized linear model for the task of model-fitting.

In the first part of the present analysis (see "Predicting Wellness Values using Workload and Sleep", below), fatigue, motivation, mood, and stress are treated as outcome variables; workload and sleep variables were treated as predictors. As outcomes, we assessed both the wellness variable and the change in that wellness variable to determine if there were effects on one or the other.

In the second part of the analysis ("Predicting Sleep Disturbance using Workload Data", below), sleep moved out of its predictor status and was treated as an outcome to assess how training might affect sleep quality. Here, practice, strength and conditioning, and competition load were used as predictors.

In the final part of the analysis ("Predicting Soreness using Workload, GPS, and Sleep", below), soreness was modeled using a larger subset of the data, involving practice, strength and conditioning, and competition load, distance, high speed distance, sprints, and sleep as predictors.

4   Results

4.1   Exploratory Data Analysis

Investigations into the dataset indicated that the position groups are meaningfully different in terms of their duties and workload. For example, figure 4.1 indicates clear differences in variables such as motivation and decelerations. The number of pairwise comparisons that would be performed to test whether each group was meaningfully different from the others would likely have created a fairly drastic Bonferroni correction and was thus not included here.

Several variables were found to be highly correlated with one another such that they could not both be included in a model together as predictors without undermining the ability to make inferences about the effect of any one of those variables on the outcome. These were as follows:

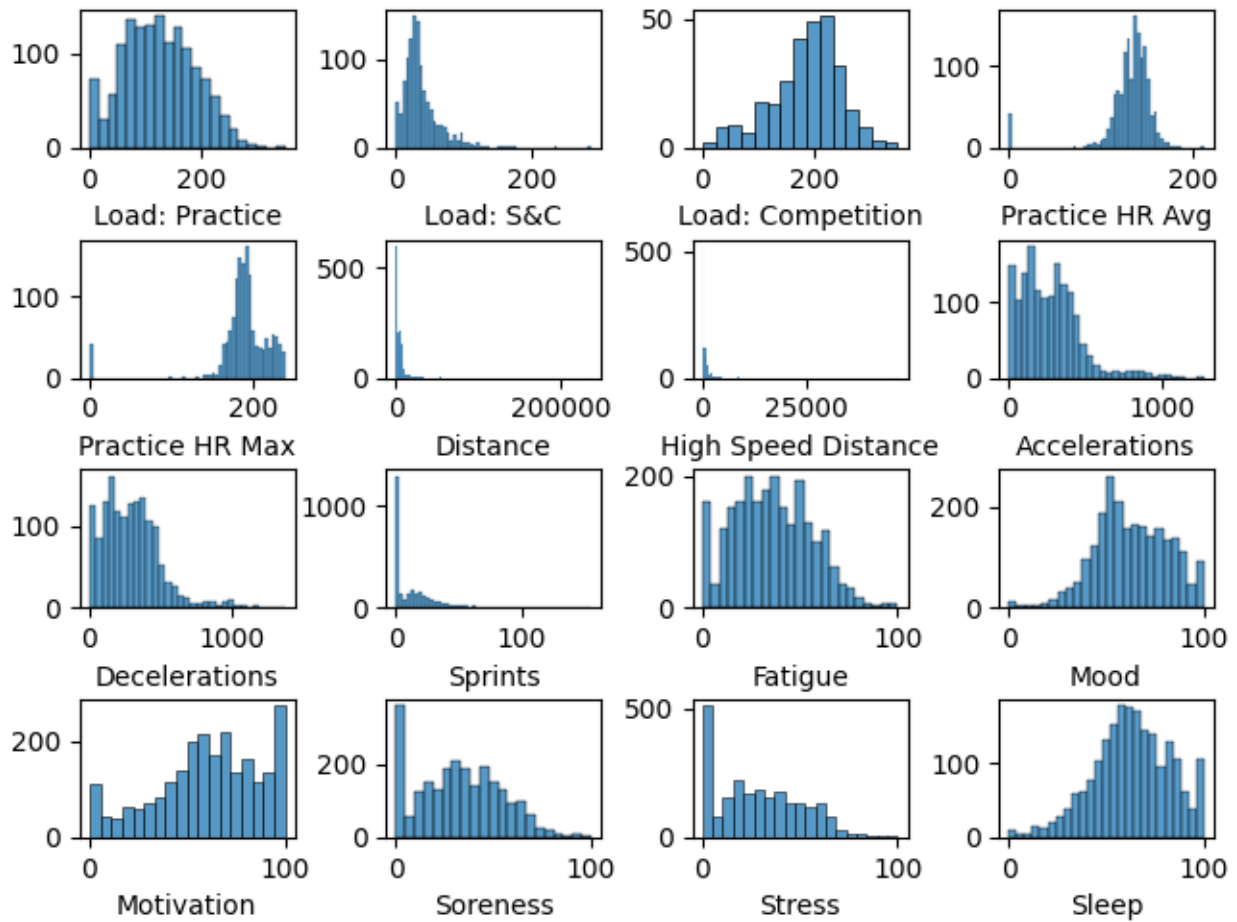- Practice load and average practice heart rate.

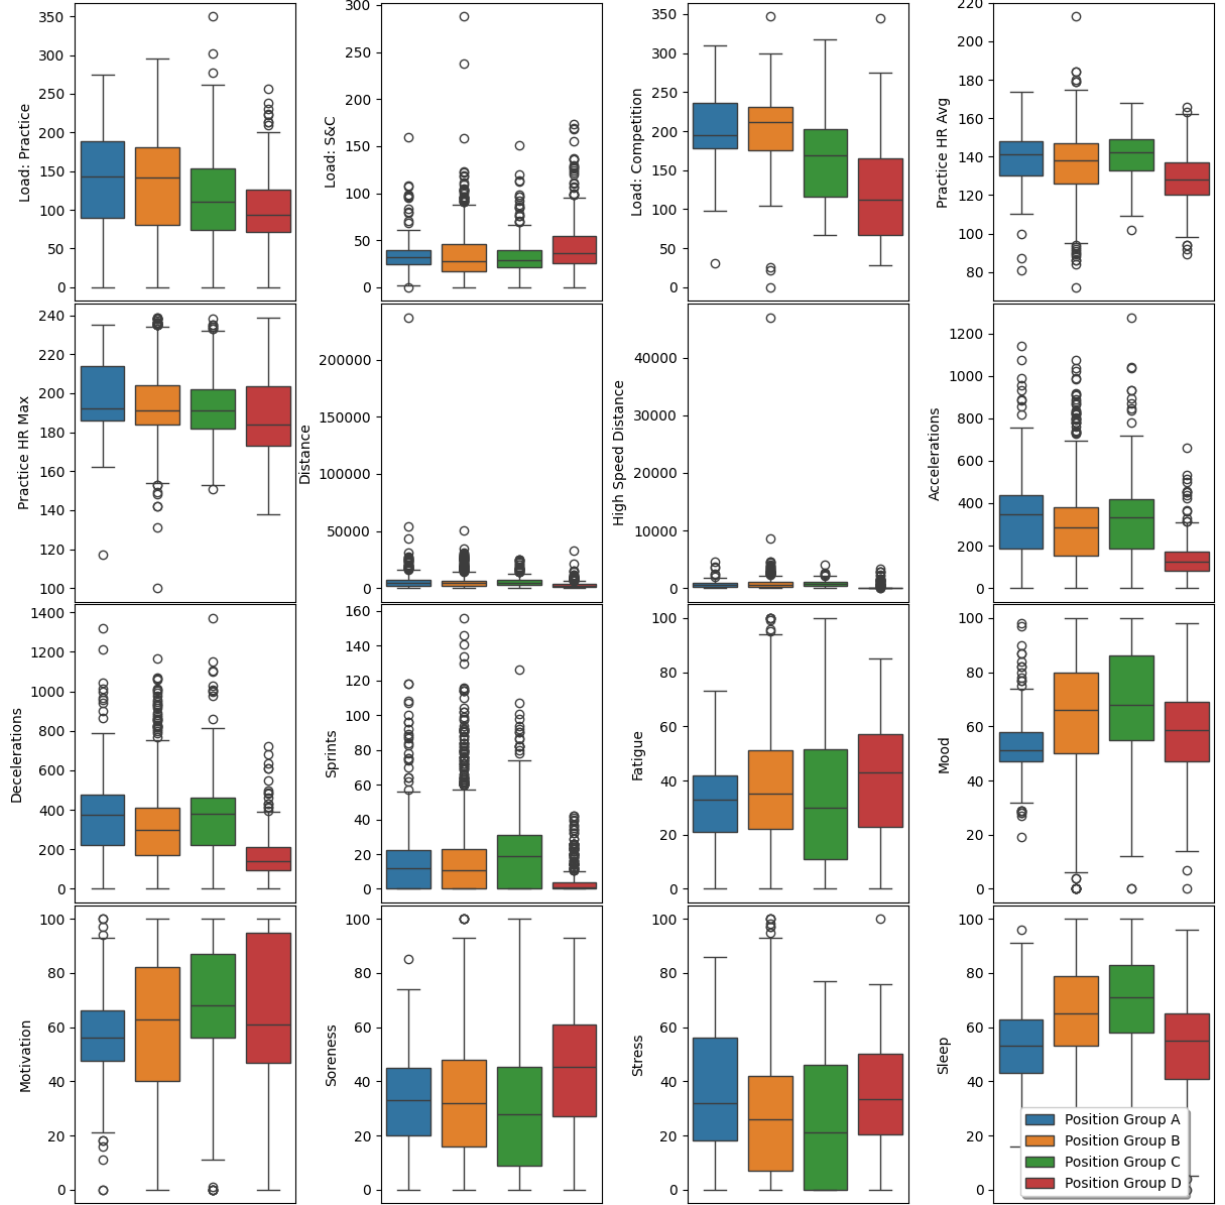Figure 1: Histograms of all numeric variables in the provided data.

Figure 2: Box and whisker plots of each variable for the four position groups. Each plot displays distributions for position groups A, B, C, and D respectively, from left to right.

- Distance and accelerations.

- Distance and decelerations.

- Accelerations and decelerations.

From this information, it was inferred that practice load and average heart rate probably capture much of the same information, which makes sense physiologically. Furthermore, simply because of the nature of GPS data, the correlations between distance and the velocity change values made sense. As a result, include practice load and heart rate were not included in the same set of predictors during modeling, nor were any two of distance, accelerations, and decelerations. Because of its overall intuitive quality and descriptiveness, distance was selected among those three as the predictor that would be included when relevant.

Exploratory data analysis also revealed that none of the wellness variables appeared to be normally distributed 1. Since the fatigue, soreness, and stress variables appeared to roughly follow a gamma distribution (or reversed gamma distribution), a gamma generalized linear model with the log link function was included among the models that were fit in the grid search for each outcome variable. The log link function was chosen due to matrix multiplication issues that occurred in `NumPy`'s and R's `glm` package code, which were not able to fit the data with some of the options they provided. This is likely due to some large values in the dataset that created consequently large matrix products. Using the log linking function circumvented this problem.

4.2   Predicting Wellness Values using Workload and Sleep

As shown in figure 3, sleep is by far the most impactful feature in terms of affecting an athlete's perceived fatigue level. Strength and conditioning workload may also be impactful, as neither of the two variables have a confidence interval overlapping with zero. Practice and competition loads have a much more modest effect. The accompanying model fit can be seen in figure 4.

In terms of change in fatigue compared to the previous week, sleep is again the most impactful variable, being associated with decreased perceptions of fatigue from one week to the next. None of the other coefficients have confidence intervals that do not overlap with zero 6.

Mood is also positively affected by sleep, with the other three coefficients having comparatively very modest effects, if any at all 8.

Concerning change in mood rating from week to week, sleep was the only variable whose coefficient confidence interval did not overlap with zero, suggesting a potential effect (see figure 9). However, the associated model fit was poor (see figure 10).

Sleep also positively affects motivation, being by far the most influential variable. The workload variables have a much lesser contribution, if any at all (see figure 11).

In the model fit for change in motivation from week to week, no coefficient had a confidence interval that did not overlap with zero (see figure 13). Additionally, the model fit was equivocal (see figure 14). Hence, there does not appear to be an effect of any workload variable or sleep on the change in motivation from one week to another.

Stress appears to be negatively associated with sleep. Although no causal claims can be made on the basis of this analysis, and because stress often anecdotally co-occurrs with lower levels of sleep, it is difficult to determine whether one variable drives the other in the present context and athlete population. Interestingly, strength and conditioning load is also negatively associated with stress (see figure 15). However, the model fit here is somewhat poor, making it difficult to justify any firm conclusions on the basis of the data (figure 16).

The model for change in stress from week to week had one of the least conclusive fits in the present analysis (see figure 16). Therefore, firm conclusions cannot be made on the basis of the model fit.

4.3   Predicting Sleep Disturbance using Workload Data

The quality of the model fit for sleep disturbance based on workload data is somewhat low due to the presence of many duplicate ground truth values. Tentatively, the null hypothesis can be retained that workload does not meaningfully affect sleep quality, with the possible exception of practice load. Although there may be

Figure 3: Coefficient magnitudes for gamma generalized linear regression of fatigue based on sleep and workload data.

Figure 4: Scatterplot of predicted and true log-transformed outcome values for gamma generalized linear regression of fatigue based on sleep and workload data.

Figure 5: Coefficient magnitudes for gamma generalized linear regression of fatigue delta based on sleep and workload data.

Figure 6: Scatterplot of predicted and true log-transformed outcome values for gamma generalized linear regression of fatigue delta based on sleep and workload data.

Figure 7: Coefficient magnitudes for gamma generalized linear regression of mood based on sleep and workload data.

Figure 8: Scatterplot of predicted and true log-transformed outcome values for gamma generalized linear regression of mood based on sleep and workload data.

Figure 9: Coefficient magnitudes for gamma generalized linear regression of mood delta based on sleep and workload data.

Figure 10: Scatterplot of predicted and true log-transformed outcome values for gamma generalized linear regression of mood delta based on sleep and workload data.

Figure 11: Coefficient magnitudes for gamma generalized linear regression of motivation based on sleep and workload data.

Figure 12: Scatterplot of predicted and true log-transformed outcome values for gamma generalized linear regression of motivation based on sleep and workload data.

Figure 13: Coefficient magnitudes for gamma generalized linear regression of motivation delta based on sleep and workload data.

Figure 14: Scatterplot of predicted and true log-transformed outcome values for gamma generalized linear regression of motivation delta based on sleep and workload data.

Figure 15: Coefficient magnitudes for gamma generalized linear regression of stress based on sleep and workload data.

Figure 16: Scatterplot of predicted and true log-transformed outcome values for gamma generalized linear regression of stress based on sleep and workload data.

Figure 17: Coefficient magnitudes for gamma generalized linear regression of stress delta based on sleep and workload data.

Figure 18: Scatterplot of predicted and true log-transformed outcome values for gamma generalized linear regression of stress delta based on sleep and workload data.

Figure 19: Coefficient magnitudes for gamma generalized linear regression of sleep based on workload data.

Figure 20: Scatterplot relating predicted sleep disturbance to true sleep disturbance.
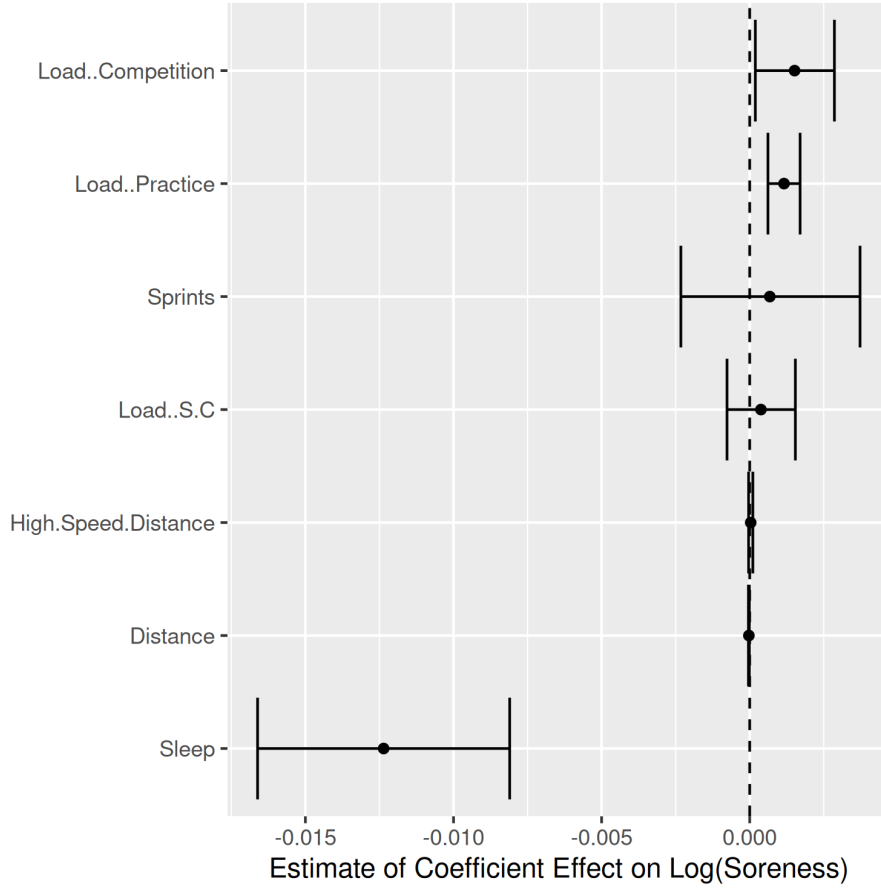
Figure 21: Coefficient magnitudes for gamma generalized linear regression of sleep based on workload data.

physiological reasons to assume that this is not true, the data do not substantiate any strong relationship between workload and sleep disturbance in this cohort and time period.

4.4 Predicting Soreness using Workload, GPS, and Sleep

In predicting soreness on the basis of sleep, load, and GPS data, sleep was once again the most powerful indicator of how sore athletes would be in a given week. Competition load may also be positively associated with soreness, having a confidence interval that did not overlap with zero (see figure 21).

5 Conclusion

5.1 Key Insights and Action Items

- EDA demonstrated that average session heart rate and practice load likely express much of the same information. If collecting both is logistically intractable or difficult, not much is lost by sacrificing one. Because of the familiarity of heart rate data and zones, it may be more useful to present these values to coaches and athletes.

- Sleep has by far the most outsized effect in promoting athlete wellness by reducing fatigue and stress and increasing mood and motivation (see figures 3 through 18).
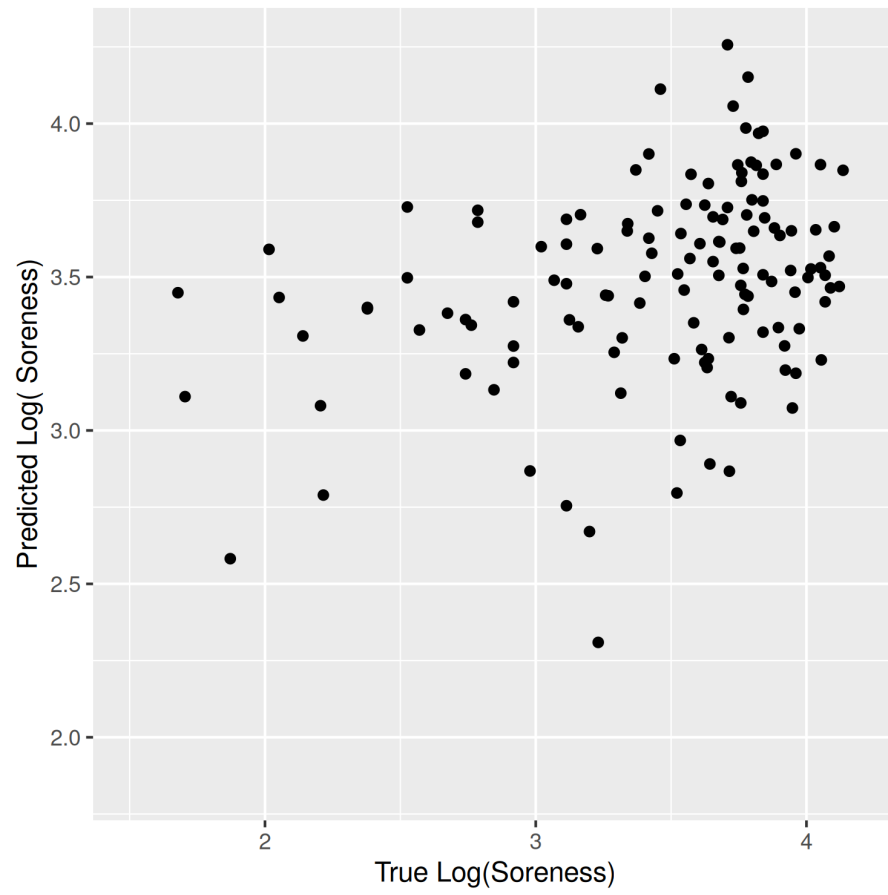
Figure 22: Scatterplot relating predicted sleep disturbance to true sleep disturbance.

- Workload does not appear to affect sleep quality within the ranges in which it has been administered (figure 19), with the exception of a potential slight effect for practice load.

- Soreness was negatively associated with sleep to a much greater degree than with any other variable (see 21). Competition load may positively influence soreness.

- There was an association between greater strength and conditioning load and lower levels of fatigue and stress (figures 3, ??). However, this effect is likely very slight and possibly confounded by position-based factors; for example, it may be the case that positions that engage in more strength and conditioning perform less other training.

- To promote athlete well-being, facilitating proper restful sleep should be considered one of the highest priorities, especially during circumstances of travel and international competition. Very secondarily, the inclusion of resistance training may help to reduce levels of fatigue and stress. Finally, since soreness is associated with high competition load, techniques to mitigate this trend may be beneficial for performance. For example, during periods of dense competition frequency when acute performance is prioritized over chronic adaptation, utilizing recovery modalities that reduce soreness may enhance performance.

5.2 Directions for Future Work

One question which this report has not addressed is the extent to which the variation in the given workload and wellness data is simply a function of athlete position group. Although there are some trends which apply across the athlete groups (e.g., the impact of sleep), and while there are certainly athletes within a particular group that might not be best served by a model of their workload and wellness that relies on their position group alone, there are probably some trends that could provide helpful coaching heuristics. This could be a fruitful direction for further investigation.

Very relatedly, it would be interesting to see if there are underlying latent positions that are different from those presented here. Principal component analysis or stochastic neighbor embedding (see [1]) could help discover if there are more underlying positions than those described by the position groups. Otremba [2] did something very similar with pitch categories in SmartPitch. Using manifold learning techniques as outlined there could provide some direction to create more tailored recommendations based on such underlying groups.

Additionally, the palette of models applied to the present task could be expanded, potentially resulted in improved outcomes. The generalized linear model with the gamma family of outcome variable distribution was chosen because of the observed density of the outcome variables included in the given data. This is valuable from an inferential standpoint, as models with assumptions that are most appropriate to the data are most suited for inferential claims about the individual coefficients. However, it may be that a Bayesian ridge regression, random forest regressor, or even basic multilayer perceptron would outperform the generalized linear model in terms of prediction accuracy. Whether or not the tradeoff in interpretability is worthwhile for these models would depend somewhat on the magnitude of the increase in prediction accuracy. However, this tradeoff could also be ameliorated to some degree with the use of Shapley values [3]. Additionally, it is questionable whether the outcomes are really gamma-distributed at all, especially given that they are expressed on a scale from 0-100. Strictly speaking, sampling from a gamma distributed variable cannot produce values equivalent to 0 and should not have a hard arbitrary upper limit the way that these outcome variables do. Hence, expanding the assumptions surrounding the distribution of the outcome variable could produce better fits and unlock further insights.

References

[1] Geoffrey E Hinton and Sam Roweis. Stochastic neighbor embedding. *Advances in neural information processing systems*, 15, 2002.

[2] Stephen Eugene Otremba Jr. *SmartPitch: Applied machine learning for professional baseball pitching strategy.* PhD thesis, Massachusetts Institute of Technology, 2022.

[3] Erik Štrumbelj and Igor Kononenko. Explaining prediction models and individual predictions with feature contributions. *Knowledge and information systems*, 41:647–665, 2014.